

Lab 1 Report

Exercise 1: Design Decisions

Tuple Class

- Attributes: Added attributes ``TupleDesc`` (instance of Tuple class), ``fields`` (to store field data), and ``recordId`` (to store tuple's disk location).
- Implementation: Followed the provided comments for function implementation, including initialization and methods to get or modify field data.

TupleDesc Class

- Attributes: Defined a ``tdItems`` list to store ``TDItem`` objects.
- Implementation: Introduced a private helper function ``initializeTdItems`` to avoid code duplication during initialization. Completed functions as per comments. Utilized Java 8 stream features to improve ``equals`` and ``getSize`` methods.

Exercise 2: Design Decisions

Table Class

- Implemented a ``Table`` class to store tables in the database.
- Attributes: Designed to include file, primary key name, and table name.
- Functions: Created methods for querying and setting table properties. Added checks for duplicate table names in the add table function.

Exercise 3: Design Decisions

BufferPool

- Implemented a hashmap to store mappings from ``PageId`` to ``Page``.
- Features: Added functionality to read from the disk if the page is not in the ``BufferPool``. Included a limit for the maximum number of pages the pool can store.

Exercise 4: Design Decisions

HeapPageId.java

- Attributes: Defined ``tableId`` (table number) and ``pageNo`` (page number).
- Implementation: Completed the code as per comments. Used the ``Object`` module for hash calculation.

RecordId.java

- Similar implementation to ``HeapPageId``.

HeapPage.java

- Calculations: Implemented formulas for calculating the number of tuples and header size.
- Slot Checking: Implemented methods to count empty slots and check if a slot is occupied by checking the corresponding bit in the bitmap.

Exercise 5: Design Decisions

Implementation

- Attributes: Defined ``f`` (file) and ``td`` (tuple) in the class.
- ``readPage``: Retrieves a specific page from the associated database file based on ``PageId``.
- Iterator: Defined ``currentPageNo`` and ``currentTupleIterator``. Implemented ``open``, ``hasNext``, ``next``, ``rewind``, and ``close`` functions for iteration. Included ``getPageTupleIterator``, a helper method to get a tuple iterator for a specified page from the ``BufferPool``.

Exercise 6: Design Decisions

Implementation

- Attributes: Defined ``tid``, ``tableId``, ``tableAlias``, and a tuple iterator ``iterator``.
- Iterator Methods: Directly utilized methods from the ``iterator`` designed in Exercise 5.

Time Spent & Challenges:

- Spent approximately 5 days on the lab.
- The most challenging part was implementing the iterator methods in Exercise 5, particularly managing the transition between different pages and ensuring the iterator correctly handles end-of-page scenarios. Also, optimizing the ``HeapPage`` calculations in Exercise 4 to efficiently compute the number of tuples and header size required a deep understanding of the underlying data structure, which was initially confusing.